OPENAI INTERNAL

Building Evaluative Trust in ChatGPT

Shifting AI from a "Black Box" to a "Glass Box"

A product proposal for Senior Product Leaders & Engineers

CHATGPT INSPECT

TRUST LAYER

THE PROBLEM

The "Fluency" Problem

Polish masks poor reasoning. Users cannot evaluate what they cannot see. The visual coherence of AI-generated text creates a false signal of accuracy – users are cognitively primed to trust output that *looks* complete.

100%

of users say outputs look complete before fully reading them

86%

of users find errors only after the fact – post-use, not pre-trust

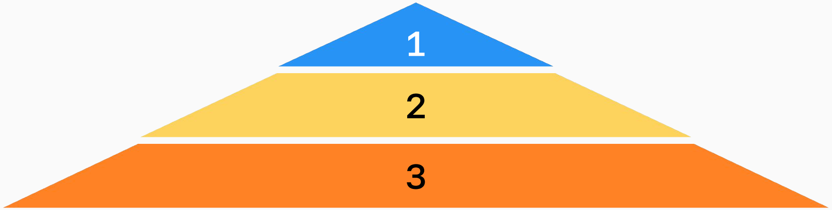

1 Errors Discovered

86% of users find mistakes only after relying on the output

2 Shallow Reading

Fluent prose triggers a "looks good" heuristic – users stop reading critically

3 Perceived Completeness

100% of users perceive outputs as complete before finishing them – the surface signals authority

BEHAVIORAL INSIGHT

The Behavioral Blind Spot

Users evaluate AI like a **search engine** – scanning for surface plausibility – rather than like an **employee** whose reasoning and assumptions must be interrogated. The result is systematic under-scrutiny of high-stakes outputs.

Critical Gap: Not a single user in research studies checked the underlying assumptions of an AI-generated response.

This is not a user failure – it is a design failure. The interface provides no affordance for assumption-checking.

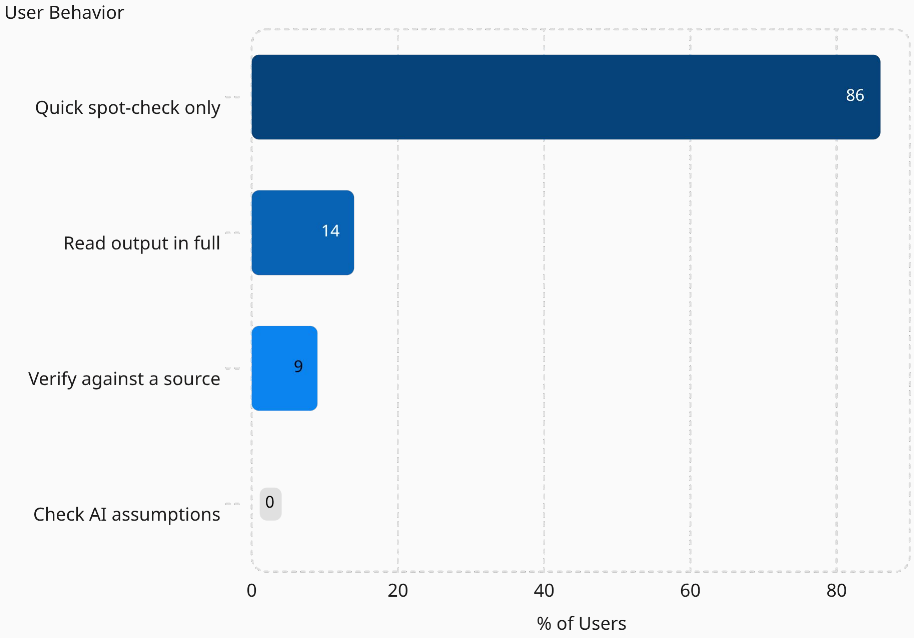

User evaluation behaviors reveal a near-total absence of assumption-checking – the most critical form of AI scrutiny.

PRODUCT VISION

The Product Vision

We must build **cognitive scaffolding** that teaches users to calibrate trust – without replacing their judgment.

🔒 AI as Black Box

Output delivered as a finished artifact. No visibility into reasoning, assumptions, or confidence. User trust is binary: accept or reject.

🔍 AI as Glass Box

Output unbundled into claims, assumptions, and alternatives. Trust becomes a **calibrated, adjustable slider** – not a leap of faith.

Key Signal: 71% of users trust their own knowledge when conflict with the AI occurs – they *want* to think critically. We just need to give them the right interface to do it.

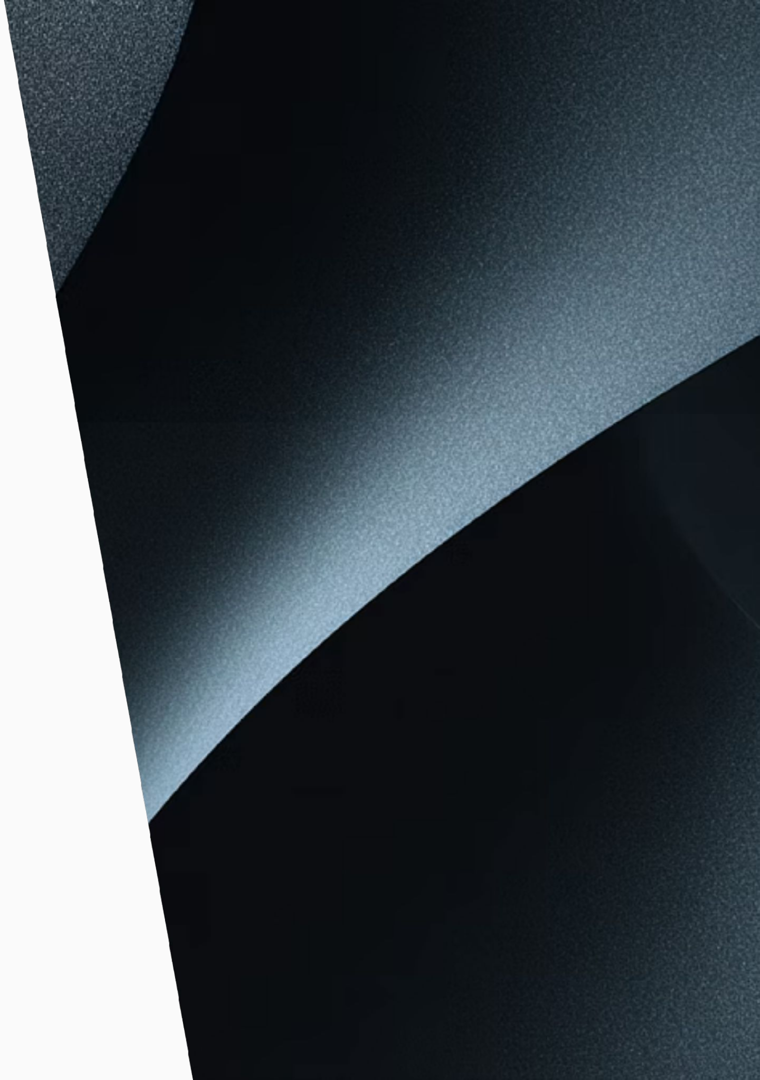

NEW FEATURE

Introducing "ChatGPT Inspect"

Unbundling the output into three evaluable layers: **Core Claims**, **Assumptions**, and **Alternative Paths** – so users can engage with AI reasoning, not just AI answers.

Core Claims

The primary assertions the AI is making, surfaced explicitly so users can challenge each one independently.

Assumptions

The hidden premises underpinning the output – directly addresses the **0% assumption-check rate** we see today.

Alternative Paths

Other valid approaches or conclusions – directly addresses the **57% of users** who want alternative viewpoints and the **43%** wanting reasoning processes.

FEATURE 1

Feature 1 - Productive Friction & Visual Anchors

Adding **intentional friction** stops the user's brain from glazing over the text – converting passive reading into active evaluation.

Visual Anchors

Falsifiable claims are highlighted inline to **optimize the 5-10 second skimming window.** Users cannot miss what must be checked.

Action-Gated Copying

The **"Copy" button is disabled** until the user interacts with at least one highlighted claim – forcing minimal engagement before content leaves the interface.

Mechanic: In a contract draft, specific legal terms are highlighted in yellow. The user must acknowledge them before the full document can be copied or exported.

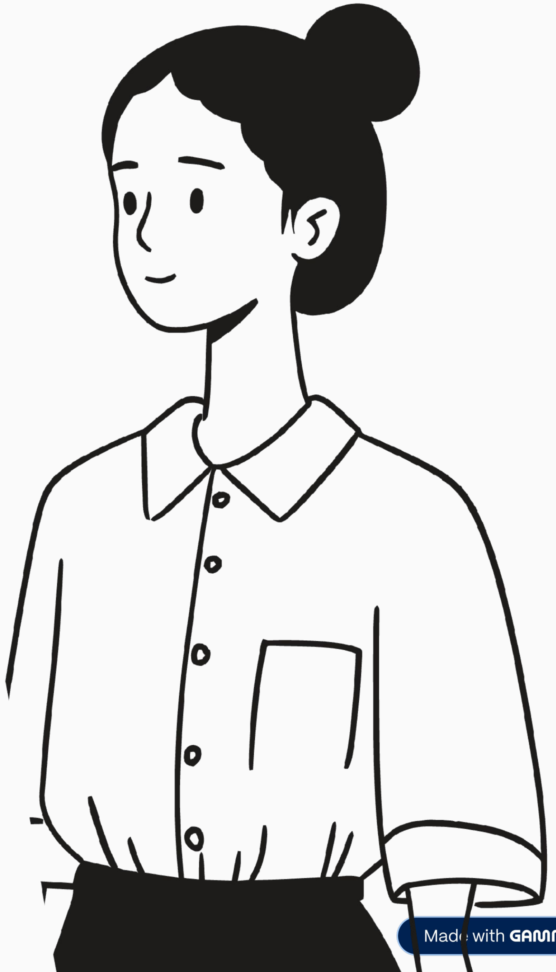

FEATURE 2

Feature 2 - Conditional Certainty

The AI explains exactly **what variable would make its output wrong** – turning trust from a binary judgment into a dynamic, user-adjustable slider.

> "If the contractor is part-time, this specific clause is legally unenforceable."

Why It Works

Leverages the **42% of users** who report trusting AI *more* when uncertainty is explicitly specified – honesty about limits builds more durable confidence than false precision.

The Trust Slider

Confidence becomes context-dependent. Users internalize the conditions of correctness, making them active participants in validation rather than passive consumers.

ARCHITECTURE

System Architecture (The "Two-Brain" System)

Using **Intent-Based Routing**, a primary AI writes the answer at full speed while a secondary "Critic AI" activates *only* for high-stakes prompts – keeping the experience fast 90% of the time.

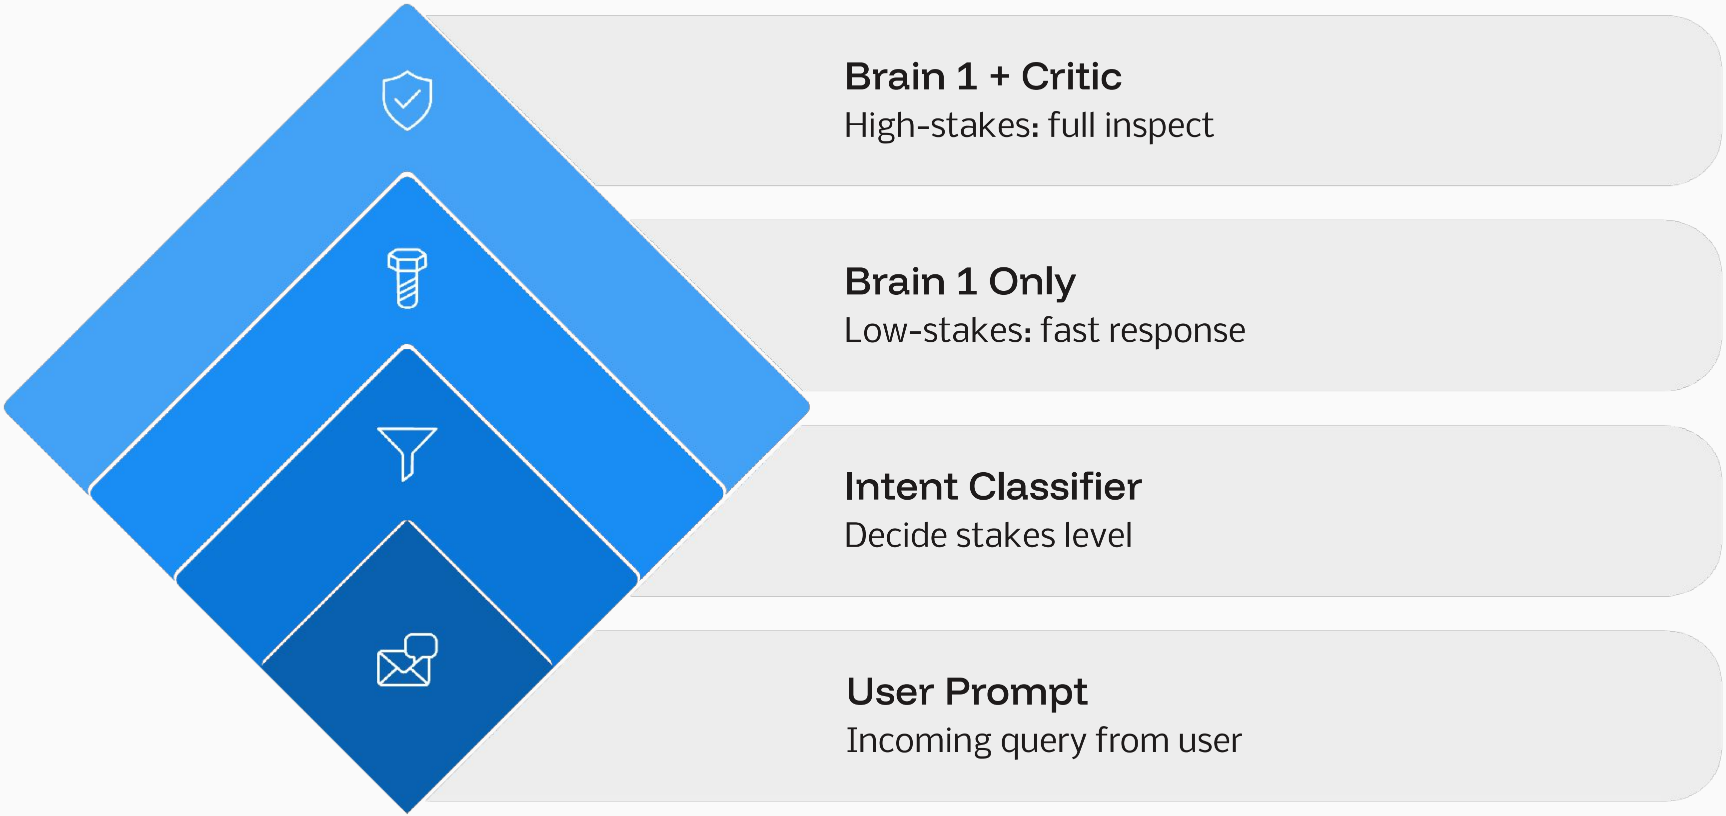

Brain 1 — Primary Model

Brain 2 — Critic AI

SUCCESS METRICS

Success Metrics

Success is measured by **active user refinement**, not passive reading time. A user who reads for 60 seconds without interacting has not demonstrated evaluative engagement.

North Star Metric

Evaluative Engagement Rate (EER): The % of high-complexity sessions where users actively interact with the Inspect panel – clicking, expanding, or dismissing a highlighted claim or assumption.

Primary: EER

% of high-complexity sessions with at least one Inspect panel interaction. This is the signal that the feature is changing behavior – not just being seen.

Secondary: Gate Completion Rate

% of users who encounter an "Action Required" gate and complete the required interaction before copying – measures friction effectiveness.

Guardrail: Trust Miscalibration Rate

% of sessions where users override a Critic AI warning and later report an error. A rising rate signals the feature is being ignored rather than used.

RISK ANALYSIS

Risk Analysis & Mitigations

Three primary risk vectors have been identified across behavioral, technical, and adoption dimensions. Each has a targeted mitigation built into the feature design.

Risk	Severity	Mitigation
Banner Blindness – Users learn to ignore the highlighted Inspect warnings the same way they ignore cookie consent banners. The feature becomes visual noise.	🔴 High	**"Action Required" Gate:** The Copy/Export button is disabled until the user interacts with at least one flagged claim. Passive ignoring is architecturally impossible – interaction is enforced at the point of highest user motivation.
Latency & Slow Load Times – Routing every prompt through a secondary Critic AI model doubles inference time, degrading the core ChatGPT experience and driving churn.	🟡 Medium	**Intent-Based Routing:** The Critic AI activates only for high-stakes, high-complexity prompts identified by the Intent Classifier. This keeps the app at standard speed for **~90% of all prompts**, limiting latency impact to a targeted minority of sessions.
Feature Abandonment – Power users who find the friction unwelcome disable or bypass the Inspect panel entirely, reducing adoption among the highest-value segment.	🟡 Medium	**Progressive Disclosure & Power Mode:** Introduce a user setting that adjusts Inspect panel intensity. Default on for all users; allow advanced users to tune sensitivity threshold without fully disabling the underlying Critic AI layer.